

An Experimental Evaluation of Containerized CI/CD Pipelines with Integrated Prometheus–Grafana Observability

Vinayak Tyagi

Department of Computer Science and Engineering
Manipal University Jaipur
vinayaktyagi.ed@gmail.com

Abstract—Containerization, CI/CD automation, and observability tooling are each widely studied on their own, but comparatively few studies measure their *combined* operational effect under a controlled experimental design. This paper builds a reproducible pipeline — Docker, GitHub Actions, Prometheus, and Grafana — around a purpose-built subject application, and experimentally evaluates it against two baselines (Baseline A: no containerization; Baseline B: Docker without CI/CD or monitoring), a third, non-monitored but CI/CD-enabled configuration (Docker+CI/CD) that isolates automation’s own cost, and against itself at increasing levels of observability once containerized (Docker only, Docker+Prometheus, Docker+Prometheus+Grafana — the last being the full proposed system). Across six controlled experiments (188 runs total: deployment time and reliability across four configurations, resource utilization and monitoring overhead across three cases, CPU-stress detection, memory-stress/OOM behaviour, and failure detection/recovery), we find that: containerized deployment adds roughly 0.86s of fixed overhead over a raw process start (1.45s vs 0.59s mean); a CI/CD-automated deployment is reliable (7/8 successful triggers) but roughly $35\times$ slower end-to-end (50.2s vs. 1.45s mean) than a pre-built local Docker deploy, because it re-pays the full build+test cost every run; a Prometheus scrape loop detects induced CPU saturation in 5.65s mean (20/20 runs) and kernel-level container failure in 1.15s mean (20/20 runs), both bounded below by the 2s scrape interval; and, counter-intuitively, the monitored application container’s own CPU utilization *decreases* as more monitoring containers are added (11.0% $\rightarrow$ 3.8% $\rightarrow$ 1.8% mean, Docker $\rightarrow$ +Prometheus $\rightarrow$ +Grafana) even as end-to-end throughput drops (2737 $\rightarrow$ 2641 $\rightarrow$ 2449 req/s), consistent with host-level CPU contention between the monitoring stack and the application rather than monitoring instrumentation overhead inside the application itself; a later replication, however, shows this CPU trend to be largely a sampling artifact and finds no measurable monitoring overhead (Section VIII). All raw per-run measurements and analysis scripts underlying these results are published at <https://github.com/vinayaktyagi10/cicd-monitoring-study>.

Index Terms—CI/CD, Docker, containerization, DevOps, Prometheus, Grafana, observability, empirical evaluation

I. INTRODUCTION

DevOps practice has converged on three largely independent pillars: *containerization* (packaging an application and its runtime into a portable, isolated unit, typically with Docker), *CI/CD automation* (build, test, and deployment triggered automatically on every change, typically via a hosted runner such as GitHub Actions), and *observability* (continuous col-

lection and visualization of runtime metrics, typically via a Prometheus scrape loop feeding a Grafana dashboard). Each pillar is well studied in isolation: containerization has an extensive performance-overhead literature going back to Felter et al. [1]; CI/CD automation is a standard subject of DevOps tooling papers; and Prometheus/Grafana integration is a common implementation detail in cloud-native systems papers. What is measured far less often is the *combined* operational profile of running all three together — specifically, how much each additional layer costs in deployment time, steady-state resource consumption, and throughput, and what it buys back in faster failure detection.

This paper reports a controlled, repeated-measures experimental evaluation of exactly that combination. We build a small subject application, wrap it in a Docker image, automate its build and test via a GitHub Actions workflow, and instrument it with Prometheus (scraping native application metrics) and Grafana (dashboarding). We then run it under six experiments spanning normal operation, CPU stress, memory stress, and induced failure, comparing it against two baselines — an uncontainerized process (Baseline A) and a containerized, CI/CD-free, unmonitored deployment (Baseline B) — and a third, intermediate configuration (Docker+CI/CD) that isolates the cost of automation alone, before monitoring is added on top to reach the full proposed system. Every run was logged in full, including two environment-specific tooling constraints encountered during the study and how each was resolved (Section VIII), so the results in Section VI are traceable back to raw measurements rather than summary claims.

A. Research Questions

- **RQ1.** How does containerization affect application deployment time and resource utilization in a CI/CD environment?
- **RQ2.** How does integrating automated monitoring affect the detection of performance degradation and failures?
- **RQ3.** What is the overhead introduced by Prometheus–Grafana monitoring in a containerized CI/CD deployment?
- **RQ4.** Does automated CI/CD combined with monitoring improve deployment reliability compared with a non-automated baseline?

B. Objectives

- 01.** Develop a reproducible containerized CI/CD pipeline using GitHub Actions.
- 02.** Integrate Prometheus and Grafana for real-time infrastructure/application monitoring.
- 03.** Measure CI/CD performance, container resource utilization, and monitoring overhead.
- 04.** Evaluate failure detection and recovery under controlled fault conditions.

II. RELATED WORK

Table I summarizes the studies most directly relevant to this work’s three pillars. This is a seed set assembled from a targeted IEEE Xplore search (queries: *Docker container performance overhead benchmark*; *containerized CI/CD pipeline performance evaluation*; *Prometheus Grafana monitoring microservices*) rather than an exhaustive systematic review, and is intended as a starting point for a full literature survey before submission.

Felter et al. [1] remains the most-cited apples-to-apples comparison of native execution, Docker containers, and KVM virtual machines, finding containers add near-zero CPU/memory overhead relative to native but that both containers and VMs pay a cost on I/O-bound workloads. Lingayat et al. [2] and Mavridis and Karatza [3] extend this line to deployment/startup time and to containers nested inside VMs respectively, but both stop at the container layer: neither introduces a CI/CD pipeline or a monitoring stack, so their measurements cannot speak to RQ2–RQ4. Ekanayaka et al. [4] is the closest system to the one built here — *KubFlow* integrates Git, Jenkins, ArgoCD, Terraform, Docker, Kubernetes, Prometheus, and Grafana into a single DevOps infrastructure — but its evaluation is a pass/fail checklist of eight qualitative test scenarios plus subjective colleague ratings, not a quantitative, repeated-measures comparison against a non-monitored or non-automated baseline.

Beyond these four, the search surfaced further candidates not yet reviewed in depth: an integrated CI/CD-pipeline study [5], a DevOps security-first CI/CD framework [6], and an AWS/Kubernetes Prometheus–Grafana deployment study [7]. These are listed for completeness and should be incorporated into a full related-work pass before this paper is submitted.

A. Research Gap

Existing studies frequently evaluate CI/CD automation, containerization, or observability independently (Table I, rows 1–3), or combine all three but report only qualitative, pass/fail integration testing rather than a controlled quantitative comparison (row 4). Comparatively few studies provide a controlled empirical evaluation of their *combined* operational impact across deployment performance, resource utilization, and failure detection, measured against explicit non-containerized and non-monitored baselines with repeated trials. This is the gap this study targets.




Fig. 1. Overall pipeline architecture: source control triggers automated build and test, which produces the container image deployed alongside the observability stack.

III. PROPOSED FRAMEWORK

Fig. 1 shows the overall pipeline architecture and data flow. A developer pushes to the Git repository, which triggers a GitHub Actions workflow that runs the test suite and, on success, builds and pushes a Docker image. The image is then run under Docker Compose alongside a cAdvisor container-metrics exporter, a Prometheus server that scrapes both the application’s native `/metrics` endpoint and cAdvisor, and a Grafana instance provisioned with Prometheus as its data source.

Fig. 2 details the GitHub Actions workflow itself, exactly as defined in the repository’s `ci.yml`: a `test` job (checkout, Python 3.12 setup, dependency install, `pytest`) gates a `docker-build` job (checkout, GHCR login, build and push) via `needs: test` — the build only runs if the tests pass. This is the real workflow triggered for every Docker+CI/CD measurement in Table IV, not a simplified stand-in.

Fig. 3 shows the resulting Grafana dashboard under live load: request rate and p95 latency by endpoint, the application’s self-reported process CPU% and memory, and the Prometheus `up` target-health indicator, all auto-provisioned (no manual dashboard configuration) and updating in real time from the same Prometheus instance used throughout the

TABLE I
RELATED WORK SUMMARY

Paper	CI/CD Docker Monitoring			Metrics	Limitation (relative to this study)
Felter et al. 2015 [1]	No	Yes (vs. KVM)	No	CPU, memory, network, disk microbenchmarks; MySQL/Redis throughput	Isolates container vs. VM overhead only; no CI/CD, no monitoring stack, no failure-detection measurement.
Lingayat et al. 2018 [2]	No	Yes (vs. bare metal / VM)	No	Container startup time	Single metric (startup time); no automation pipeline, no monitoring, no fault injection.
Mavridis & Karatza 2017 [3]	No	Yes (nested in VM)	No	Standard container benchmarks	Studies an orthogonal deployment topology (container-in-VM); no CI/CD or monitoring integration.
Ekanayaka et al. 2023 [4]	Yes (Jenkins+ArgoCD)	Yes (Kubernetes)	Yes (Prometheus+Grafana)	Pass/fail test scenarios; subjective 1–5 feature ratings	Same three pillars as this study, but evaluation is qualitative/subjective, not a controlled quantitative comparison against a non-monitored or non-automated baseline.
This study	Yes (GitHub Actions)	Yes (Docker Compose)	Yes (Prometheus+Grafana)	Deployment time, reliability, CPU/mem utilization, throughput, latency, MTTD, MTTR — $n=20$ per configuration where feasible	—




Fig. 2. GitHub Actions workflow (.github/workflows/ci.yml). The docker-build job only runs if test succeeds (needs: test); this is the real workflow triggered for the Docker+CI/CD deployment-time measurements in Table IV.

experiments.

Subject application. Rather than an existing multi-service demo application, the study uses a small, hand-built FastAPI service with three endpoints designed as clean experimental levers: /health (a near-constant-time route used as the load-test and liveness target), /cpu (a configurable-iteration pure-Python busy loop used for CPU stress), and /memory




Fig. 3. Grafana dashboard under live load, auto-provisioned against the Prometheus datasource. The request-rate and CPU panels show a live traffic spike generated during the capture; the target-health panel (bottom, not shown in this crop) reports up=1.

(a configurable-size heap allocation, optionally retained, used for memory stress). This choice trades away the multi-service realism of a demo application such as Sock Shop for the ability to attribute every measured effect to a specific, known cause rather than to inter-service scheduling or network confounds. The application exposes Prometheus metrics natively via the prometheus_client library — a request counter and latency histogram per endpoint, plus two self-reported process gauges, app_process_cpu_percent and app_process_memory_bytes, sampled every second from time.process_time() and /proc/self/status deltas respectively — rather than through a side-car exporter.

IV. EXPERIMENTAL SETUP

Table II lists the experimental environment. All experiments ran on a single physical host (no cloud VM); GitHub Actions

TABLE II
EXPERIMENTAL ENVIRONMENT

Parameter	Value
CPU	AMD Ryzen 7 8845HS (8 cores / 16 threads)
RAM	14 GiB
Host OS	Arch Linux (Omarchy), kernel 7.1.9-arch1-2
Docker	29.7.2 (containerd snapshotter storage driver)
GitHub Actions	ubuntu-latest hosted runners
Prometheus	v2.55.1, 2s scrape/evaluation interval
Grafana	11.2.0
cAdvisor	v0.49.1
Application base image	python:3.12-slim
Application framework	FastAPI 0.115.0 / Uvicorn 0.30.6
Load generator	Dependency-free Python ThreadPoolExecutor + urllib script (no external binary installed)

TABLE III
EXPERIMENTAL SCENARIOS

#	Scenario	Configurations × runs
1	Deployment time & reliability	manual (20), Docker (20), full pipeline (20), Docker+CI/CD (8)
2	Resource utilization, response time, throughput, monitoring overhead	Docker (20), Docker+Prometheus (20), Docker+Prometheus+Grafana (20)
3	CPU stress detection	20
4	Memory stress / OOM behaviour	20
5 & 6	Failure detection & recovery	20

hosted runners were used only for the CI/CD build-and-test step itself, so that step reflects a genuinely independent, reproducible environment while deployment, load generation, and monitoring remained under full local control.

A load-testing tool such as `hey` or `wrk` was deliberately not installed; a small standard-library Python script (`loadgen.py`) issues concurrent HTTP requests via `ThreadPoolExecutor` and records per-request latency, trading raw throughput ceiling for zero external dependencies and full inspectability.

V. EXPERIMENTAL SCENARIOS

Six experiments were run, summarized in Table III. Each is described in more detail in the corresponding results subsection. A dedicated script per experiment (`exp1_deploy.py` through `exp5_failure_recovery.py`) brings the relevant Docker Compose profile up, executes the scenario, records raw per-run measurements to CSV, and tears the stack down before the next run, so runs are independent of one another. An `aggregate.py` script then computes the mean/standard deviation/min/max summaries reported in Section VI directly from these raw CSV files.

A. Normal operation and baseline comparison (Experiment 1)

Four configurations are compared, only two of which are baselines proper: *manual* (a raw `uvicorn` process, no container — **Baseline A**) and *Docker* (a locally built and started

container, no CI/CD, no monitoring — **Baseline B**). The remaining two are not baselines but an explicit progression toward the proposed system: *Docker+CI/CD* (a real GitHub Actions `workflow_dispatch` run — test job then Docker build/push — timed end-to-end, then deployed) is an **intermediate configuration** that isolates what automation alone costs, with monitoring still absent; and *full pipeline* (Docker Compose bringing up the application alongside `cAdvisor`, `Prometheus`, and `Grafana`) is the **proposed system**, adding observability on top. This ordering — Baseline A → Baseline B → Docker+CI/CD → full pipeline — lets each of containerization, automation, and observability be attributed its own marginal cost rather than only the combined effect of all three at once. For each run, whatever was previously running is torn down, the configuration is started, and wall-clock time to the first successful health check (and, for the full pipeline, to `Prometheus` and `Grafana` also responding) is recorded.

B. Resource utilization and monitoring overhead (Experiment 2)

Three cases isolate the marginal cost of each monitoring layer: Docker only, Docker+Prometheus(+`cAdvisor`), and Docker+Prometheus+Grafana(+`cAdvisor`). For each of 20 runs per case, a fixed load (300 concurrent requests to `/health`, concurrency 10) is issued while `docker stats` is polled every second for every container in that case, giving per-container mean CPU% and memory alongside the load test’s own throughput, mean/p95 latency, and error rate.

C. CPU stress detection (Experiment 3)

Eight threads concurrently flood `/cpu` (an 8-million-iteration pure-Python loop per call) against the Docker+Prometheus case. Detection time is measured as wall-clock time from load start until `Prometheus` first reports the application’s self-reported `app_process_cpu_percent` gauge at or above 60%, polled every 0.5s.

D. Memory stress / OOM behaviour (Experiment 4)

The application container is run directly (not via `Compose`) with a hard 150 MB memory limit and `--restart on-failure:1`. Repeated `/memory?mb=20&hold=true` calls accumulate retained allocations until the kernel OOM-kills the container; `docker inspect` gives ground-truth `OOMKilled` and `RestartCount`, captured at the moment of observation since Docker resets `State.OOMKilled` once the restart policy brings the container back up.

E. Application failure and recovery (Experiments 5 & 6)

Against the Docker+Prometheus case, the application container is killed (`docker kill`) and `Prometheus`’s `up{job="app"}` series is polled every 0.3s until it reports 0 (detection time). The container is then restarted (`docker start`) and both `/health` and `up{job="app"}==1` are polled until both recover (recovery time).

TABLE IV
DEPLOYMENT TIME BY CONFIGURATION (SECONDS)

Config	Mean	Std	Min	Max	N (fail)
Manual (Baseline A)	0.592	0.020	0.577	0.633	20 (0)
Docker (Baseline B)	1.454	0.107	1.267	1.581	20 (0)
Docker+CI/CD	50.231	3.817	44.404	55.362	8 (1)
Full pipeline	1.552	0.031	1.485	1.636	20 (0)




Fig. 4. Deployment time by configuration (log scale; error bars are one standard deviation). Docker adds a small fixed cost over the raw process; the full monitoring stack adds almost nothing further on top of Docker; CI/CD dominates because it re-runs the full build-and-test cycle rather than reusing a pre-built image.

VI. RESULTS

A. Deployment time and reliability (RQ1, RQ4)

Table IV reports wall-clock deployment time for the three fast configurations and CI/CD-inclusive deployment time for the fourth. Containerizing a fixed local image adds a consistent ≈ 0.86 s of fixed overhead over a raw process start (1.454s vs. 0.592s mean); adding the full monitoring stack on top of Docker adds a further ≈ 0.10 s (1.552s vs. 1.454s mean) since Compose brings up four containers rather than one but all four start in parallel. The CI/CD-automated path is roughly $35\times$ slower (50.2s vs. 1.454s mean) because it re-pays the full test-then-build cost on every deployment rather than reusing a pre-built image — this is the direct cost of automation, traded against the reliability and repeatability of not requiring a human to run the equivalent commands by hand. Of 8 triggered CI/CD runs, 7 completed successfully; the one failure was traced to a transient delay in the workflow being registered by GitHub’s API immediately after being triggered, rather than to a defect in the pipeline itself, and is reported as observed rather than discarded.

Fig. 4 visualizes Table IV on a log scale, since the CI/CD configuration is roughly 30–85 $\times$ slower than the other three and would otherwise flatten them to indistinguishable bars.

TABLE V
APPLICATION RESOURCE USE AND RESPONSE CHARACTERISTICS

Config	App CPU%	App Mem (MB)	p95 (ms)	Throughput (req/s)
Docker	11.02 (3.23)	36.11 (0.41)	5.23 (1.00)	2737 (206)
+Prometheus	3.76 (4.42)	35.68 (0.43)	5.34 (0.80)	2641 (226)
+Grafana	1.79 (3.11)	35.79 (0.96)	5.51 (0.86)	2449 (285)

Mean (std. dev.), $n = 20$ runs of 300 requests each per row.

B. Resource utilization, throughput, and monitoring overhead (RQ1, RQ3)

Table V reports the application container’s own resource use and the load test’s response characteristics; Table VI isolates the monitoring stack’s own resource draw. Zero errors were observed across all 60 runs ($180\times 300 = 18,000$ requests). Point estimates for throughput fall as monitoring is added ($2737 \rightarrow 2641 \rightarrow 2449$ req/s, Docker $\rightarrow$ +Prometheus $\rightarrow$ +Grafana); a Welch’s t -test (unequal variances, $n=20$ per case, computed directly from the raw per-run CSV rather than from the summary means) confirms the two endpoints of this trend are statistically distinguishable (Docker vs. +Grafana: $t=3.57$, $p=0.0011$, Cohen’s $d=1.13$), but the middle step is not (Docker vs. +Prometheus: $t=1.37$, $p=0.177$) — the overall direction is real, but the drop is not uniform run-to-run between every adjacent pair. Mean and p95 latency both rise numerically but *none* of the three pairwise differences reach significance at $\alpha=0.05$ (all $p \geq 0.36$); this claim is therefore stated only as a directional observation, not a confirmed effect, and should not be read as strong evidence on its own. Counter-intuitively, the application container’s own mean CPU% *falls* over the same progression ($11.0\% \rightarrow 3.8\% \rightarrow 1.8\%$), and here the statistical evidence is strong: Docker vs. +Prometheus is highly significant ($t=5.78$, $p<0.0001$, $d=1.83$, a large effect), though +Prometheus vs. +Grafana is not ($p=0.12$) — the CPU-share cost of adding Prometheus itself is real and large, while adding Grafana on top is not distinguishable from noise at this sample size. As cAdvisor, Prometheus, and Grafana compete for the same physical cores, the application container is scheduled less rather than working harder, so its own `docker stats` CPU reading drops even as end-to-end throughput degrades. This contention reading is a hypothesis, not an established finding: the load window is only ≈ 0.11 s, and the `docker stats` reading covers a longer interval, so many runs read as idle (0/20, 12/20 and 16/20 runs below 1% CPU for Docker, +Prometheus and +Grafana) while non-idle runs read ≈ 8 –11% in every case; a follow-up measurement from kernel counters did not support it (Section VIII). The `prometheus_client` counters/histograms/gauges add negligible per-request work, as expected.

Table VI is visualized in Fig. 10, alongside the failure-detection and MTTD/MTTR figures at the end of this section.

C. CPU stress detection (RQ2)

All 20 induced CPU-stress runs were detected by Prometheus, with mean time to detect (MTTD) 5.65s (std 3.68s, min 1.51s, max 18.58s) and mean peak self-reported




Fig. 5. Application container CPU utilization across the three monitoring cases (Table V, column 1). CPU% *falls* as more monitoring containers are added — the counter-intuitive result discussed under RQ3.




Fig. 6. Application container memory across the three monitoring cases (Table V, column 2; note the narrow y-axis range). Memory is essentially flat at ≈ 36 MB regardless of how much monitoring runs alongside the application — unlike CPU, memory is not shared/contended the same way across containers on this host.

CPU utilization 96.6% (std 10.0, effectively pegged at one full core, consistent with CPython’s GIL limiting a single process to $\approx 100\%$ of one core regardless of thread count). The detection floor is set by the 2s Prometheus scrape interval combined with the 0.5s poll granularity used to observe it; the long right tail (up to 18.58s) reflects run-to-run variance in how quickly the OS scheduler favoured the eight stress threads over the sampling thread, not a monitoring-pipeline delay.

An initial design used mean `/health` latency as the detection signal instead of process CPU%; it was rejected after finding it barely moved under load, because CPython’s GIL round-robins between threads at a short, fixed interval (≈ 5 ms) regardless of how CPU-bound the competing work is, so a trivial handler’s queueing delay stays small even under heavy contention.




Fig. 7. Mean and p95 request latency across the three monitoring cases (Table V, columns 3-4). Both rise slightly as monitoring is added, consistent with the throughput drop and the host-contention explanation in Section VI.

TABLE VI
MONITORING STACK’S OWN RESOURCE DRAW (TABLE 5: OVERHEAD)

Container	+Prometheus case		+Grafana case	
	CPU%	Mem (MB)	CPU%	Mem (MB)
cAdvisor	6.13	47.16	6.40	45.39
Prometheus	0.78	59.40	0.80	52.78
Grafana	—	—	0.06	60.70

Mean, $n = 20$ runs per column.

D. Memory stress and OOM behaviour (RQ2)

All 20 runs triggered a genuine kernel OOM-kill and automatic restart (`RestartCount=1`, confirmed via `docker inspect`), after a mean 5.25 successful allocation requests (std 1.09) and a mean peak reported memory of 135.54 MB (std 0.17) against the 150 MB hard limit. Mean time-to-OOM was 11.58s (std 2.23s); this figure is dominated by the fixed per-iteration overhead of the polling loop itself (a `curl`-equivalent request, a `docker inspect` call, and a `docker stats` call per iteration) rather than by the kernel’s actual OOM response latency, and should not be read as a monitoring detection time in the same sense as Experiments 3, 5, and 6.

E. Failure detection and recovery (RQ2, RQ4)

All 20 induced container failures (`docker kill`) were detected by Prometheus’s `up{job="app"}` series, mean time to detect (MTTD) 1.15s (std 0.21s, min 0.82s, max 2.03s) — close to the theoretical floor set by the 2s scrape interval combined with 0.3s poll granularity and Docker’s own signal-delivery latency. Mean time to recover (MTTR) — container restart to health-check success — was consistently ≈ 2.0 s (std 0.006s, i.e. effectively deterministic), which is essentially Uvicorn’s own process startup time inside the container: MTTR here is dominated by application boot, not by any property of the monitoring stack. From this point on, MTTD and MTTR are used throughout in place of “detection time” and “recovery time,” consistent with the CPU-stress MTTD reported above and with Figs. 8–9.




Fig. 8. MTDD and MTTR across all 20 failure-injection runs. MTDD (blue) is consistently faster and tighter than MTTR (orange) across every run, and both are stable run-to-run — neither shows a trend over the sequence of 20 trials, supporting the claim that these are steady-state properties of the pipeline rather than artifacts of run order.




Fig. 9. Mean time to detect (MTDD) a CPU-stress event, MTDD for a killed container, and mean time to recover (MTTR) after restart, with error bars (one standard deviation), computed from the same raw detection/recovery times reported in Section VI. Binary failure (container killed) is detected roughly $5\times$ faster than gradual CPU saturation, since a scrape either observes the target down or it doesn't, while a saturation threshold must be crossed by an accumulating signal.

F. Same-host replication

Experiments 1–5 were repeated with the unmodified scripts on the same machine on a later date (20 runs per configuration), writing to a separate directory so the data behind Tables III–VI is untouched. Table VII gives the original and replication means with 95% confidence intervals (`validate.py`); a metric is marked *consistent* only if the original mean lies inside the replication's interval. Most absolute values moved well outside the original intervals, while the relative structure that the research questions rely on held for deployment ordering, MTTR and peak CPU under stress (Section VIII lists what did not reproduce and why). The original run remains the paper's reported dataset because it was the pre-planned measurement and every table is derived




Fig. 10. Monitoring stack containers' own CPU draw (Table VI). cAdvisor dominates the monitoring stack's own overhead at $\approx 6\%$ CPU in both cases; Prometheus and Grafana are comparatively cheap ($<1\%$ each). This is the overhead the monitoring *tooling* itself pays, separate from the host-contention effect on the application shown in Fig. 5.

from it; the replication is reported as a robustness check, not as a replacement.

VII. DISCUSSION

Taken together, these results answer the four research questions as follows.

RQ1 (containerization's effect on deployment time and resource use). Containerization has a small, consistent fixed cost at deployment time ($\approx 0.86s$ here) and a resource footprint dominated by the application's own working set (36MB) rather than by container overhead per se — consistent with the classic finding in Felter et al. [1] that containers add near-zero CPU/memory overhead relative to native execution. The more consequential RQ1 finding in this study is that a CI/CD-automated deployment is not simply "Docker plus a trigger": it is roughly $35\times$ slower end-to-end because it re-runs the full build-and-test cycle, a cost invisible to studies (such as Lingayat et al. [2]) that measure only container startup time in isolation.

RQ2 (monitoring's effect on detecting degradation and failure). Prometheus reliably (20/20 in both experiments) detects both an induced CPU saturation event (mean 5.65s) and an induced process failure (mean 1.15s), with detection latency in both cases bounded below by the 2s scrape interval. The roughly $5\times$ difference between them (5.65s vs. 1.15s) is itself informative: a binary up/down signal is detected almost as fast as the scrape loop allows, while a gradual resource-saturation signal takes longer because it must accumulate enough samples above threshold amid scheduling noise — a distinction a purely qualitative evaluation such as Ekanayaka et al.'s pass/fail test scenarios [4] cannot surface.

RQ3 (monitoring overhead). In the original run the overhead appeared structurally different from what is typically reported: the monitoring stack's own containers (principally cAdvisor at $\approx 6\%$ CPU) appeared to compete with the application for the same fixed core budget, degrading end-to-

TABLE VII

ORIGINAL VS. SAME-HOST REPLICATION (20 RUNS EACH). INTERVALS ARE t -BASED 95% CIS ON THE MEAN (SAMPLE SD); $\Delta\%$ IS REPLICATION RELATIVE TO ORIGINAL. THE EXPERIMENT 4 ROWS ARE FROM THE UNMODIFIED ZERO-FILLED SCRIPT (0/20 OOM KILLS, SEE SECTION VIII); THE NON-ZERO-FILL VARIANT REPRODUCED THE ORIGINAL. GENERATED FROM `RESULTS/ORIGINAL_VS_REPLICATION.CSV`.

Metric	Original mean [95% CI]	Replication mean [95% CI]	$\Delta\%$	Verdict
Deploy time, manual (s)	0.59 [0.58, 0.60]	0.42 [0.41, 0.43]	-28.9	different
Deploy time, Docker (s)	1.45 [1.40, 1.50]	1.05 [1.04, 1.06]	-27.7	different
Deploy time, full pipeline (s)	1.55 [1.54, 1.57]	1.37 [1.31, 1.43]	-11.8	different
Deploy time, Docker+CI/CD (s)	—	—	—	not run
Throughput, Docker (req/s)	2737.07 [2638.29, 2835.84]	4040.66 [3984.28, 4097.04]	+47.6	different
Throughput, +Prometheus (req/s)	2640.65 [2532.08, 2749.23]	3995.78 [3944.85, 4046.70]	+51.3	different
Throughput, +Grafana (req/s)	2448.85 [2311.85, 2585.84]	4033.22 [3995.02, 4071.43]	+64.7	different
p95 latency, Docker (ms)	5.23 [4.75, 5.70]	3.16 [2.96, 3.36]	-39.5	different
p95 latency, +Prometheus (ms)	5.34 [4.95, 5.72]	3.14 [3.03, 3.25]	-41.2	different
p95 latency, +Grafana (ms)	5.51 [5.09, 5.92]	3.03 [2.97, 3.08]	-45.1	different
App CPU%, Docker	11.02 [9.47, 12.56]	7.42 [6.37, 8.46]	-32.7	different
App CPU%, +Prometheus	3.76 [1.64, 5.88]	3.50 [1.81, 5.18]	-7.1	consistent
App CPU%, +Grafana	1.79 [0.29, 3.28]	1.91 [0.65, 3.16]	+6.6	consistent
App mem, Docker (MB)	36.11 [35.91, 36.31]	37.66 [37.49, 37.82]	+4.3	different
App mem, +Prometheus (MB)	35.68 [35.47, 35.88]	37.74 [37.50, 37.98]	+5.8	different
App mem, +Grafana (MB)	35.79 [35.32, 36.25]	37.72 [37.49, 37.95]	+5.4	different
cAdvisor CPU%, +Prometheus	6.12 [4.58, 7.67]	2.41 [1.80, 3.02]	-60.7	different
cAdvisor CPU%, +Grafana	6.40 [4.98, 7.82]	2.19 [1.60, 2.78]	-65.8	different
Prometheus CPU%, +Prometheus	0.78 [0.40, 1.16]	0.62 [0.26, 0.98]	-20.4	consistent
Prometheus CPU%, +Grafana	0.80 [0.47, 1.13]	0.51 [0.18, 0.84]	-36.1	consistent
Grafana CPU%, +Grafana	0.06 [0.01, 0.11]	0.05 [0.03, 0.07]	-22.5	consistent
CPU-stress MTTD (s)	5.65 [3.88, 7.42]	2.41 [2.07, 2.75]	-57.3	different
CPU-stress peak CPU%	96.59 [91.79, 101.39]	95.13 [89.69, 100.57]	-1.5	consistent
Time to OOM (s)	—	—	—	not run
Peak memory (MB)	135.54 [135.46, 135.62]	138.59 [137.53, 139.65]	+2.3	different
Requests before OOM	5.25 [4.73, 5.77]	16.00 [16.00, 16.00]	+204.8	different
Failure MTTD (s)	1.15 [1.05, 1.25]	3.14 [3.07, 3.20]	+172.9	different
MTTR (s)	1.99 [1.99, 2.00]	1.98 [1.94, 2.02]	-0.7	consistent

end throughput (-10.5% Docker→+Grafana) while the application’s own measured CPU% fell. A later controlled re-measurement (Section VIII) did not confirm this: the monitoring containers drew ≈ 0.03 cores in total and no throughput contrast against Docker-only was distinguishable from zero, so the contention framing should be treated as unconfirmed.

RQ4 (does automated CI/CD plus monitoring improve reliability). The evidence here is mixed and should be stated honestly: CI/CD automation was reliable (7/8, 87.5%) but not perfect, and the one observed failure was attributable to a transient delay in the trigger step rather than to the pipeline’s own robustness. Critically, 87.5% is a point estimate that should not be quoted on its own: the 95% Wilson score interval for 7 successes out of 8 trials is [52.9%, 97.8%] — at this sample size the data are consistent with anywhere from roughly half to nearly all CI/CD deployments succeeding, and the point estimate alone overstates the precision of the claim. A larger n for the CI/CD arm (constrained here by real GitHub Actions minutes and wall-clock cost, Section VIII) would be needed to narrow this interval and bound the reliability claim more tightly.

VIII. LIMITATIONS

- **Replication on the same host did not reproduce most absolute numbers.** Re-running Experiments 1–5 with the unmodified scripts on the same machine on a later date (20 runs each, new kernel, performance power profile) (Table VII) gave deployment times 12–29% faster, throughput about 48% higher, CPU-stress detection of 2.41s rather than 5.65s and failure detection

of 3.14s rather than 1.15s, with the mean outside the original 95% interval for most metrics. Only MTTR, peak CPU under stress, the Prometheus/Grafana CPU figures and the application CPU% in the +Prometheus and +Grafana cases stayed consistent. The deployment ordering (manual, then Docker, then full pipeline) was preserved, but the throughput ordering across monitoring cases was not (4041, 3996 and 4033 req/s for Docker, +Prometheus and +Grafana); absolute values are specific to one host state on one day.

- **The CPU-contention explanation is not supported.** Experiment 2’s application CPU% is bimodal because its ≈ 0.11 s load window is sampled by a longer `docker stats` interval, so the reported fall with monitoring is mostly runs that read as idle. A new Experiment 6 (20 runs per case, fixed 20s load, CPU from kernel counters over exactly the load window, randomized case order, optional CPU-pinned arm) found the monitoring containers used ≈ 0.03 cores in total and no throughput difference from Docker-only whose 95% interval excluded zero. The pinning test is inconclusive (difference-in-differences +376 req/s, 95% CI [-343, +1061]). Run-to-run throughput varied roughly three-fold and tracked host load ($r = -0.98$ with host busy cores), so the experiment can rule out large monitoring overhead on this host but not small effects.
- **Experiment 4 depends on kernel behaviour.** The retained allocation is zero-filled. On the replication kernel, which enables transparent hugepage shrinking of zero-

filled pages (`shrink_underused=1`), the kernel reclaimed that memory under cgroup pressure, so 0/20 containers were OOM-killed even though 16 requests (≈ 320 MB) succeeded against a 150 MB limit. With the allocation writing a non-zero byte to every page (`/memory?fill=true`, default off), 20/20 runs were OOM-killed after a mean 5.0 requests with peak memory 136.4 MB, matching the original result (5.25 requests, 135.5 MB). The OOM finding therefore holds when the memory is genuinely resident; the original zero-filled variant is kernel dependent. We did not establish whether the original 2026-09 run differed from the replication in THP settings or kernel.

- **cAdvisor container-label resolution.** This host’s Docker uses the containerd-snapshotter storage driver rather than classic overlay2, which cAdvisor v0.49.1 cannot resolve container name/image labels against. CPU-stress and failure-detection experiments therefore use the application’s own native Prometheus metrics rather than cAdvisor’s per-container series; Table VI’s per-container figures come from `docker stats` directly. This is an environment/tooling limitation rather than a consequence of the experimental design.
- **GHCR authentication.** The Docker+CI/CD deployment step could not pull from GitHub Container Registry (the CLI token used lacks the `read:packages` scope, which would require an interactive browser re-authentication out of scope for an unattended experiment run) and instead deploys the locally built image from the same commit and Dockerfile CI just built. The CI timing itself (test + build + push) is unaffected and fully real.
- **Reduced n for the CI/CD arm.** Only 8 (vs. 20 for the other three deployment configurations) real GitHub Actions runs were triggered, because each run costs ≈ 45 – 65 s of real CI wall time and consumes Actions minutes; this narrows the confidence interval on the CI/CD reliability and timing figures relative to the other three configurations.
- **Single-host measurement.** All experiments ran on one physical machine with no isolation from other host activity beyond what Docker itself provides; absolute timing figures (especially the sub-2s deployment times) should be read as relative comparisons between configurations on this host, not as portable absolute benchmarks.
- **Related work is a seed set.** Section II’s four fully-reviewed papers and three title-only candidates were found via a small number of targeted IEEE Xplore queries, not a systematic multi-database review; a submission-ready version of this paper should expand this section substantially, including non-IEEE venues (ACM, Springer, USENIX) before finalizing the research-gap claim.
- **Detection thresholds are fixed, not tuned.** The 60% CPU-percent threshold (Experiment 3) and the memory container limit (150 MB, Experiment 4) were chosen to make the induced condition reliably observable within

a short polling window rather than to model a specific production SLO; detection-time figures should be read relative to these thresholds, not as threshold-independent constants.

- **Not every reported difference is statistically significant.** Welch’s t -tests on the raw per-run data (`statistical_analysis.py`) confirm the Docker-to-Grafana endpoints of the throughput and CPU% trends in Section VI ($p < 0.01$ for both), but two things reported only as directional observations do not reach significance at $\alpha = 0.05$: the latency increase across all three monitoring cases ($p \geq 0.36$ for every pair), and the middle step of the throughput trend (Docker vs. +Prometheus, $p = 0.177$). Means and standard deviations alone, without this check, would have overstated how uniformly these effects hold run-to-run.

IX. CONCLUSION

This study built a small but complete containerized CI/CD pipeline with integrated Prometheus–Grafana observability and evaluated it experimentally against a non-containerized baseline and against itself at increasing monitoring depth, across six controlled experiments totaling 188 runs. The headline findings are that containerization’s own deployment-time and resource cost is small and consistent with prior literature, that CI/CD automation’s real cost is the repeated build-test cycle rather than the deployment step itself, that Prometheus reliably detects both induced resource saturation and induced failure within low-single-digit seconds bounded by its scrape interval, and that the dominant cost of running a monitoring stack alongside the monitored application on constrained hardware is host-level CPU contention between them rather than in-process instrumentation overhead — a conclusion the replication reported in Section VIII did not confirm. The full raw measurement set, pipeline configuration, and analysis scripts underlying every table in this paper are published at <https://github.com/vinayaktyagi10/cicd-monitoring-study> to support independent verification.

ACKNOWLEDGMENT

The author thanks Dr. Dibakar Sinha for his guidance and supervision throughout this study.

REFERENCES

- [1] W. Felter, A. Ferreira, R. Rajamony, and J. Rubio, “An updated performance comparison of virtual machines and Linux containers,” in *2015 IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*. Philadelphia, PA, USA: IEEE, 2015, pp. 171–172.
- [2] A. Lingayat, R. R. Badre, and A. K. Gupta, “Performance evaluation for deploying docker containers on baremetal and virtual machine,” in *2018 3rd International Conference on Communication and Electronics Systems (ICCES)*. Coimbatore, India: IEEE, 2018.
- [3] I. Mavridis and H. Karatza, “Performance and overhead study of containers running on top of virtual machines,” in *2017 IEEE 19th Conference on Business Informatics (CBI)*. Thessaloniki, Greece: IEEE, 2017.
- [4] E. M. I. M. Ekanayaka, J. K. K. H. Thathsarani, D. S. Karunanayaka, N. Kuruwitaarachchi, and N. Skandhakumar, “Enhancing DevOps infrastructure for efficient management of microservice applications,” in *2023 IEEE International Conference on e-Business Engineering (ICEBE)*. Sydney, Australia: IEEE, 2023.

- [5] S. Kumar, N. Bala, A. P. Singh, and Y. Raj, "Cloud-native continuous integration/continuous deployment (CI/CD) pipeline," in *2024 Second International Conference on Advanced Computing & Communication Technologies (ICACCTech)*. IEEE, 2024, title, authorship and venue verified via IEEE Xplore search; full text not reviewed at drafting time.
- [6] R. Ashtagi, C. Belani, P. Bhalerao, Y. Bhalerao, and A. Chawle, "Building resilient CICD pipelines: A DevOps security-first framework," in *2025 International Conference on Computational, Communication and Information Technology (ICCCIT)*. IEEE, 2025, title, authorship and venue verified via IEEE Xplore search; full text not reviewed at drafting time.
- [7] T. Abirami, S. Mapari, P. Jayadharshini, L. Krishnasamy, and R. R. Vigneshwaran, "Streamlined deployment and monitoring of cloud-native applications on AWS with Kubernetes Prometheus Grafana," in *2023 International Conference on Advances in Computation, Communication and Information Technology (ICAICCTIT)*. IEEE, 2023, title, authorship and venue verified via IEEE Xplore search; full text not reviewed at drafting time.